

Solving Triangular Systems

- 1 Forward Substitution Formulas
 - processing the L of the LU factorization
 - a third type of pipeline
- 2 Multiple Double Arithmetic
 - ill conditioned matrices
 - quad double arithmetic
- 3 Parallel Solving
 - using an n -stage pipeline
 - rewriting the formulas
 - a parallel solver with OpenMP

MCS 572 Lecture 29
Introduction to Supercomputing
Jan Verschelde, 19 March 2021

Solving Triangular Systems

1 Forward Substitution Formulas

- processing the L of the LU factorization
- a third type of pipeline

2 Multiple Double Arithmetic

- ill conditioned matrices
- quad double arithmetic

3 Parallel Solving

- using an n -stage pipeline
- rewriting the formulas
- a parallel solver with OpenMP

LU factorization

To solve an n -dimensional linear system $A\mathbf{x} = \mathbf{b}$
we factor A as a product of two triangular matrices, $A = LU$:

- L is lower triangular, $L = [\ell_{i,j}]$, $\ell_{i,j} = 0$ if $j > i$ and $\ell_{i,i} = 1$.
- U is upper triangular $U = [u_{i,j}]$, $u_{i,j} = 0$ if $i > j$.

Solving $A\mathbf{x} = \mathbf{b}$ is equivalent to solving $L(U\mathbf{x}) = \mathbf{b}$:

- 1 Forward substitution: $L\mathbf{y} = \mathbf{b}$.
- 2 Backward substitution: $U\mathbf{x} = \mathbf{y}$.

Factoring A costs $O(n^3)$, solving triangular systems costs $O(n^2)$.

formulas for forward substitution

Expanding the matrix-vector product $L\mathbf{y}$ in $L\mathbf{y} = \mathbf{b}$ leads to

$$\begin{cases} y_1 & = b_1 \\ \ell_{2,1}y_1 + y_2 & = b_2 \\ \ell_{3,1}y_1 + \ell_{3,2}y_2 + y_3 & = b_3 \\ & \vdots \\ \ell_{n,1}y_1 + \ell_{n,2}y_2 + \ell_{n,3}y_3 + \cdots + \ell_{n,n-1}y_{n-1} + y_n & = b_n \end{cases}$$

and solving for the diagonal elements gives

$$\begin{aligned} y_1 &= b_1 \\ y_2 &= b_2 - \ell_{2,1}y_1 \\ y_3 &= b_3 - \ell_{3,1}y_1 - \ell_{3,2}y_2 \\ &\vdots \\ y_n &= b_n - \ell_{n,1}y_1 - \ell_{n,2}y_2 - \cdots - \ell_{n,n-1}y_{n-1} \end{aligned}$$

formula and algorithm

For $k = 1, 2, \dots, n$:

$$y_k = b_k - \sum_{i=1}^{k-1} \ell_{k,i} y_i.$$

As an algorithm:

for k from 1 to n do

$y_k := b_k$;

 for i from 1 to $k - 1$ do

$y_k := y_k - \ell_{k,i} \star y_i$.

We count

$$1 + 2 + \dots + n - 1 = \frac{n(n-1)}{2}$$

multiplications and subtractions.

Solving Triangular Systems

1 Forward Substitution Formulas

- processing the L of the LU factorization
- a third type of pipeline

2 Multiple Double Arithmetic

- ill conditioned matrices
- quad double arithmetic

3 Parallel Solving

- using an n -stage pipeline
- rewriting the formulas
- a parallel solver with OpenMP


a third type of pipeline

Three types of pipelines:

- 1 Speedup only if multiple instances. Example: instruction pipeline.
- 2 Speedup already if one instance. Example: pipeline sorting.
- 3 Worker continues after passing information through.

Example: solve $L\mathbf{y} = \mathbf{b}$.

Typical for the 3rd type of pipeline is the varying length of each job.



Solving Triangular Systems

1 Forward Substitution Formulas

- processing the L of the LU factorization
- a third type of pipeline

2 Multiple Double Arithmetic

- ill conditioned matrices
- quad double arithmetic

3 Parallel Solving

- using an n -stage pipeline
- rewriting the formulas
- a parallel solver with OpenMP

ill conditioned matrices

Consider the 4-by-4 lower triangular matrix

$$L = \begin{bmatrix} 1 & 0 & 0 & 0 \\ -2 & 1 & 0 & 0 \\ -2 & -2 & 1 & 0 \\ -2 & -2 & -2 & 1 \end{bmatrix}.$$

What we know from numerical analysis:

- 1 The condition number of a matrix magnifies roundoff errors.
- 2 The hardware double precision is $2^{-52} \approx 2.2 \times 10^{-16}$.
- 3 We get no accuracy from condition numbers larger than 10^{16} .

an experiment in an interactive Julia session

```
julia> using LinearAlgebra

julia> A = ones(32,32);

julia> D = Diagonal(A);

julia> L = LowerTriangular(A);

julia> LmD = L - D;

julia> L2 = D - 2*LmD;

julia> cond(L2)
2.41631630569077e16
```

The condition number is estimated at 2.4×10^{16} .

Solving Triangular Systems

1 Forward Substitution Formulas

- processing the L of the LU factorization
- a third type of pipeline

2 Multiple Double Arithmetic

- ill conditioned matrices
- quad double arithmetic

3 Parallel Solving

- using an n -stage pipeline
- rewriting the formulas
- a parallel solver with OpenMP

quad double arithmetic

A quad double is an unevaluated sum of 4 doubles, improves working precision from 2.2×10^{-16} to 2.4×10^{-63} .

Y. Hida, X.S. Li, and D.H. Bailey: **Algorithms for quad-double precision floating point arithmetic**. In *15th IEEE Symposium on Computer Arithmetic* pages 155–162. IEEE, 2001. Software at

<http://crd.lbl.gov/~dhbailey/mpdist/qd-2.3.17.tar.gz>.

A quad double builds on `double double`, some features:

- The least significant part of a `double double` can be interpreted as a compensation for the roundoff error.
- Predictable overhead: working with `double double` is of the same cost as working with complex numbers.

operator overloading in C++

```
#include <iostream>
#include <iomanip>
#include <qd/qd_real.h>
using namespace std;

int main ( void )
{
    qd_real q("2");
    cout << setprecision(64) << q << endl;
    for(int i=0; i<8; i++)
    {
        qd_real dq = (q*q - 2.0)/(2.0*q);
        q = q - dq; cout << q << endl;
    }
    cout << scientific << setprecision(4);
    cout << "residual : " << q*q - 2.0 << endl;
    return 0;
}
```

compiling with a makefile

The makefile contains the entry:

```
QD_ROOT=/usr/local/qd-2.3.17
```

```
QD_LIB=/usr/local/lib
```

```
qd4sqrt2:
```

```
    g++ -I$(QD_ROOT)/include qd4sqrt2.cpp \  
        $(QD_LIB)/libqd.a \  
        -o qd4sqrt2
```

Compiling at the command prompt \$:

```
$ make qd4sqrt2
```

```
g++ -I/usr/local/qd-2.3.17/include qd4sqrt2.cpp \  
    /usr/local/lib/libqd.a \  
    -o qd4sqrt2
```

running the code

```
$ ./qd4sqrt2
2.00000000000000000000000000000000000000000000000000000000000e+00
1.50000000000000000000000000000000000000000000000000000000000e+00
1.41666666666666666666666666666666666666666666666666666666667e+00
1.4142156862745098039215686274509803921568627450980392156862745098e+00
1.4142135623746899106262955788901349101165596221157440445849050192e+00
1.4142135623730950488016896235025302436149819257761974284982894987e+00
1.4142135623730950488016887242096980785696718753772340015610131332e+00
1.4142135623730950488016887242096980785696718753769480731766797380e+00
1.4142135623730950488016887242096980785696718753769480731766797380e+00
residual : 0.0000e+00
$
```

General multiple double arithmetic is available:

M. Joldes, J.-M. Muller, V. Popescu, W. Tucker.

CAMPARY: Cuda Multiple Precision Arithmetic Library and Applications.

In *Mathematical Software – ICMS 2016, the 5th International Conference on Mathematical Software*, pages 232–240, Springer-Verlag, 2016.

Solving Triangular Systems

1 Forward Substitution Formulas

- processing the L of the LU factorization
- a third type of pipeline

2 Multiple Double Arithmetic

- ill conditioned matrices
- quad double arithmetic

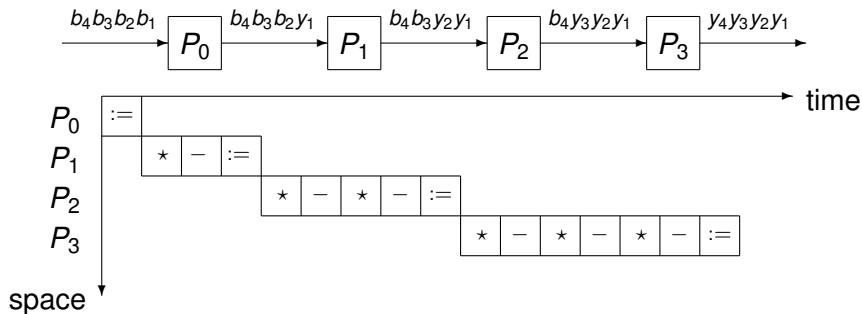
3 Parallel Solving

- using an n -stage pipeline
- rewriting the formulas
- a parallel solver with OpenMP

using an n -stage pipeline


We assume that L is available on every processor.

for $n = 4 = p$: $y_1 := b_1$
 $y_2 := b_2 - \ell_{2,1} \star y_1$
 $y_3 := b_3 - \ell_{3,1} \star y_1 - \ell_{3,2} \star y_2$
 $y_4 := b_4 - \ell_{4,1} \star y_1 - \ell_{4,2} \star y_2 - \ell_{4,3} \star y_3$




type 3 pipelining

Make y_1 available in the next pipeline cycle:



counting the steps



We count the steps for $p = 4$ or in general, for $p = n$:

- 1 The latency takes 4 steps for y_1 to be at P_4 ,
or in general: n steps for y_1 to be at P_n .
- 2 It takes then 6 additional steps for y_4 to be computed by P_4 ,
or in general: $2n - 2$ additional steps for y_n to be computed by P_n .

So it takes $n + 2n - 2 = 3n - 2$ steps to solve
an n -dimensional triangular system by an n -stage pipeline.

Solving Triangular Systems

1 Forward Substitution Formulas

- processing the L of the LU factorization
- a third type of pipeline

2 Multiple Double Arithmetic

- ill conditioned matrices
- quad double arithmetic

3 Parallel Solving

- using an n -stage pipeline
- **rewriting the formulas**
- a parallel solver with OpenMP

rewriting the formulas

Solving $L\mathbf{y} = \mathbf{b}$ for $n = 5$:

① $y := \mathbf{b};$

② $y_2 := y_2 - \ell_{2,1} \star y_1;$

$$y_3 := y_3 - \ell_{3,1} \star y_1;$$

$$y_4 := y_4 - \ell_{4,1} \star y_1;$$

$$y_5 := y_5 - \ell_{5,1} \star y_1;$$

③ $y_3 := y_3 - \ell_{3,2} \star y_2;$

$$y_4 := y_4 - \ell_{4,2} \star y_2;$$

$$y_5 := y_5 - \ell_{5,2} \star y_2;$$

④ $y_4 := y_4 - \ell_{4,3} \star y_3;$

$$y_5 := y_5 - \ell_{5,3} \star y_3;$$

⑤ $y_5 := y_5 - \ell_{5,4} \star y_4;$

$$\mathbf{y} := \mathbf{b};$$

for i from 2 to n do

for j from i to n do

$$y_j := y_j - \ell_{j,i-1} \star y_{i-1};$$

\Rightarrow all instructions in the j loop are independent from each other!

data parallelism

Consider the inner loop in the algorithm to solve $L\mathbf{y} = \mathbf{b}$:

$\mathbf{y} := \mathbf{b}$;

for i from 2 to n do

for j from i to n do

$y_j := y_j - \ell_{j,i-1} \star y_{i-1}$;

We distribute the update of y_i, y_{i+1}, \dots, y_n among p processors.

If $n \gg p$, then we expect a close to optimal speedup.

Solving Triangular Systems

1 Forward Substitution Formulas

- processing the L of the LU factorization
- a third type of pipeline

2 Multiple Double Arithmetic

- ill conditioned matrices
- quad double arithmetic

3 Parallel Solving

- using an n -stage pipeline
- rewriting the formulas
- a parallel solver with OpenMP

a parallel solver

For our parallel solver for triangular systems:

- For $L = [\ell_{i,j}]$, we generate random numbers for $\ell_{i,j} \in [0, 1]$.
The exact solution \mathbf{y} : $y_i = 1$, for $i = 1, 2, \dots, n$.
We compute the right hand side $\mathbf{b} = L\mathbf{y}$.
- Even already in small dimensions,
the condition number may grow exponentially.
Hardware double precision is insufficient.
Therefore, we use quad double arithmetic.
- We use a straightforward OpenMP implementation.

solving random lower triangular systems

Relying on hardware doubles is problematic:

```
$ time ./trisol 10
```

```
last number : 1.00000000000000009e+00
```

```
real    0m0.003s    user    0m0.001s    sys     0m0.002s
```

```
$ time ./trisol 100
```

```
last number : 9.9999999999974221e-01
```

```
real    0m0.005s    user    0m0.001s    sys     0m0.002s
```

```
$ time ./trisol 1000
```

```
last number : 2.7244600009080568e+04
```

```
real    0m0.036s    user    0m0.025s    sys     0m0.009s
```

a matrix of quad doubles

Allocating data in the main program:

```
{  
    qd_real b[n], y[n];  
  
    qd_real **L;  
    L = (qd_real**) calloc(n, sizeof(qd_real*));  
    for(int i=0; i<n; i++)  
        L[i] = (qd_real*) calloc(n, sizeof(qd_real));  
  
    srand(time(NULL));  
    random_triangular_system(n, L, b);  
}
```

a random triangular system

```
void random_triangular_system
( int n, qd_real **L, qd_real *b )
{
    for(int i=0; i<n; i++)
    {
        L[i][i] = 1.0;
        for(j=0; j<i; j++)
        {
            double r = ((double) rand())/RAND_MAX;
            L[i][j] = qd_real(r);
        }
        for(int j=i+1; j<n; j++)
            L[i][j] = qd_real(0.0);
    }
    for(int i=0; i<n; i++)
    {
        b[i] = qd_real(0.0);
        for(int j=0; j<n; j++)
            b[i] = b[i] + L[i][j];
    }
}
```

solving the system

```
void solve_triangular_system_swapped
( int n, qd_real **L, qd_real *b, qd_real *y )
{
    for(int i=0; i<n; i++) y[i] = b[i];

    for(int i=1; i<n; i++)
    {
        for(int j=i; j<n; j++)
            y[j] = y[j] - L[j][i-1]*y[i-1];
    }
}
```

using OpenMP

```
void solve_triangular_system_swapped
( int n, qd_real **L, qd_real *b, qd_real *y )
{
    int j;

    for(int i=0; i<n; i++) y[i] = b[i];

    for(int i=1; i<n; i++)
    {
        #pragma omp parallel shared(L,y) private(j)
        {
            #pragma omp for
            for(j=i; j<n; j++)
                y[j] = y[j] - L[j][i-1]*y[i-1];
        }
    }
}
```

experimental timings

running time `./trisol_qd_omp n p`

On dimension $n = 8,000$ for varying number p of cores.

p	cpu time	real	user	sys
1	21.240s	35.095s	34.493s	0.597s
2	22.790s	25.237s	36.001s	0.620s
4	22.330s	19.433s	35.539s	0.633s
8	23.200s	16.726s	36.398s	0.611s
12	23.260s	15.781s	36.457s	0.626s

The serial part is the generation of the random numbers for L and the computation of $\mathbf{b} = L\mathbf{y}$. Recall Amdahl's Law.

We can compute the serial time, subtracting for $p = 1$, from the real time the cpu time spent in the solver, i.e.: $35.095 - 21.240 = 13.855$.
For $p = 12$, time spent on the solver is $15.781 - 13.855 = 1.926$.
Compare 1.926 to $21.240/12 = 1.770$.

Summary + Exercises

We ended chapter 5 in the book of Wilkinson and Allen.

Exercises:

- 1 Consider the upper triangular system $U\mathbf{x} = \mathbf{y}$, with $U = [u_{i,j}]$, $u_{i,j} = 0$ if $i > j$. Derive the formulas and general algorithm to compute the components of the solution \mathbf{x} .
For $n = 4$, draw the third type of pipeline.
- 2 Write a parallel solver with OpenMP to solve $U\mathbf{x} = \mathbf{y}$.
Take for U a matrix with random numbers in $[0, 1]$, compute \mathbf{y} so all components of \mathbf{x} equal one. Test the speedup of your program, for large enough values of n and a varying number of cores.
- 3 Describe a parallel solver for upper triangular systems $U\mathbf{y} = \mathbf{b}$ for distributed memory computers. Write a prototype implementation using MPI and discuss its scalability.